

THE UNIVERSITY OF CHICAGO

REMARKABLE FEATURES: USING DESCRIPTIVE CONTRAST TO EXPRESS AND
INFER TYPICALITY

A DISSERTATION SUBMITTED TO
THE FACULTY OF THE DIVISION OF THE SOCIAL SCIENCES
IN CANDIDACY FOR THE DEGREE OF
DOCTOR OF PHILOSOPHY

DEPARTMENT OF PSYCHOLOGY

BY
CLAIRE AUGUSTA BERGEY

CHICAGO, ILLINOIS
AUGUST 2023

TABLE OF CONTENTS

LIST OF FIGURES	iv
LIST OF TABLES	vi
ACKNOWLEDGMENTS	vii
ABSTRACT	ix
INTRODUCTION	1
1 PEOPLE TALK MORE ABOUT ATYPICAL THAN TYPICAL FEATURES OF THINGS	6
1.1 Adjective typicality	8
1.1.1 Corpora	9
1.1.2 Stimulus Selection	9
1.1.3 Participants	10
1.1.4 Design and Procedure	11
1.1.5 Results	11
1.1.6 Discussion	15
1.2 Extracting Typicality from Language Structure	16
1.2.1 Method	17
1.2.2 Results	18
1.3 General Discussion	21
2 HOW ADULTS USE CONTRASTIVE INFERENCE TO LEARN ABOUT NEW CATEGORIES	24
2.1 Experiment 1	26
2.1.1 Method	26
2.1.2 Results	29
2.1.3 Discussion	31
2.1.4 Model	32
2.2 Experiment 2	36
2.2.1 Method	36
2.2.2 Results	38
2.2.3 Model	40
2.2.4 Discussion	41
2.3 Conclusion	42
3 HOW CHILDREN USE CONTRASTIVE INFERENCE TO LEARN ABOUT NEW CATEGORIES	43
3.1 Method	45
3.1.1 Participants.	45

3.1.2	Design and Procedure.	47
3.2	Performance on practice trials	48
3.3	Results	48
3.4	Discussion	51
CONCLUSION		52
REFERENCES		56

LIST OF FIGURES

1.1	Density plots showing parents' use of atypical and typical adjective-noun pairs across their child's age.	12
1.2	Density plots showing use of atypical and typical adjective-noun pairs by parents speaking to children and adults speaking to other adults.	13
1.3	Proportion of caregiver description that is about highly typical features (often, almost always, or always true), as a function of age.	14
1.4	Density plots showing children's use of atypical and typical adjective-noun pairs across age.	15
1.5	Plots of word2vec noun-adjective similarities for nouns for which there was at least one atypical adjective (rated at most "sometimes"), and at least one typical adjective (rated at least "often").	20
1.6	Plots of BERT and GPT-3 noun-adjective similarities for nouns for which there was at least one atypical adjective (rated at most "sometimes"), and at least one typical adjective (rated at least "often").	20
2.1	Experiment 1 stimuli. In the above example, the critical feature is size and the object context is a within-category contrast: the alien on the right has two same-shaped objects that differ in size.	26
2.2	Participants' prevalence judgments from Experiment 1, along with our model predictions. Participants consistently judged the target object as less typical of its category when the referent was described with an adjective (e.g., "Pass me the blue toma") than when it was not (e.g., "Pass me the toma"). This inference was not significantly modulated by object context (examples shown above each figure panel). Points indicate empirical means; error bars indicate 95% confidence intervals computed by non-parametric bootstrapping. Solid horizontal lines indicate model predictions.	35
2.3	Participants' prevalence judgments in Experiment 2, with model predictions using the parameters estimated in Experiment 1 (horizontal lines).	40
3.1	An example of the novel objects shown on a trial. In each trial, two objects of the same shape and differing on the critical feature were shown sequentially. In adjective noun trials, the critical feature was mentioned for the object that had it (e.g., the wide toma was called a "wide toma") and in noun trials, no features were mentioned (e.g., both tomas were just called a "toma".	46
3.2	An example of the prevalence judgment children were asked to make. Children chose between clouds of novel objects representing few, some, most, and almost all of the novel category having the feature. The experimenter asked, e.g., "Let's think about all of the blickets on this planet. How many blickets do you think are spotted? Few of the blickets, some of the blickets, most of the blickets, or almost all of the blickets?"	46
3.3	Children's prevalence judgments across utterance conditions and feature types. .	49

3.4	Prevalence judgments among only children who answered the practice trials correctly. These children rate features to be less prevalent when they are mentioned with an adjective.	50
-----	---	----

LIST OF TABLES

- 1.1 Sample typicality ratings from three human coders for three adjective-noun pairs drawn from the corpus. Note that means may be slightly different from the mean of the three ratings shown here because some pairs have more than three ratings. 10

ACKNOWLEDGMENTS

Each chapter in this dissertation represents collaborative work. The collaborators on each chapter are as follows: Chapter 1, Benjamin Morris and Dan Yurovsky; Chapter 2, Dan Yurovsky; Chapter 3, Jordyn Martin. None of this work could have happened without Ben, Dan, and Jordyn, and anything this dissertation gets right can be credited to their intellectual contributions to these chapters and to my life.

Thank you to my committee, Susan Goldin-Meadow, Marisa Casillas, Howard Nusbaum, and Dan Yurovsky for all they have given me: frameworks to think in, an intellectual community to be a part of, and guidance to follow (or not). There are few gifts better than a new way to see the world.

I am deeply grateful for people who have given feedback on this work along the way: Ming Xiang, Benjamin Morris, Ashley Leung, Michael C. Frank, Judith Degen, Stephan Meylan, Robert Hawkins, and Ruthe Foushee. Thanks also to reviewers and attendees of Experiments in Linguistic Meaning, the meeting of the Cognitive Science Society, and the Midwestern Cognitive Science Conference, and the journal *Cognition*.

Many people made up my life while I wrote this—too many to mention. I’m pretty sure I can’t think without talking to people, and you all are my best interlocutors. A few choice words for choice people:

To Jean and Barry, for the love of thinking and art—you made me and made me who I am; To Matthew, it’s a pleasure being siblings with you, wouldn’t have it any other way;

To Julia, for being my best critic—seeing what the best of something can be; To Rosemary, relentless collaborator, you help me keep hold of my self;

To Scott, who has seen it all happen, for talking with me all this time; To Russell, here’s to us; To Jane and Erica, my oldest friends;

To Ben, for tackling our research and our friendship; To Hannah, for being strange and remembering why we like each other; To Yağmur, for challenging and care in equal measure;

To Casey, for learning to teach together, and for the jokes; To Ashley, we did it; To Jimmy, you kept me from taking it all too seriously;

To Ruthe, for telling me you can do whatever you want and being right; To Ivan, for bringing people together; To Mike, for arguing with me and sitting with me.

I love you all. None of you will ever get rid of me.

ABSTRACT

We mention what is remarkable while letting the unremarkable go unsaid. Thus, while language can tell us a lot about the world, it does not veridically reflect the world: people are more likely to talk about atypical features (e.g., “purple carrot”) than typical features (e.g., “[orange] carrot”). In this dissertation, I characterize how people selectively describe the features of things and examine the implications of this selective description for how children and adults learn from language. In Chapter 1, I show that adults speaking to other adults, caregivers speaking to children, and children themselves tend to mention the atypical more than the typical features of concrete things. Language is structured to emphasize what is atypical—so how can one learn about what things are typically like from language? In this chapter I also show that language models that use word co-occurrence to derive word meaning do not capture the typicality of adjective–noun pairs well. However, people can learn in more complex ways: in Chapter 2, I show that people infer that mentioned features are atypical. That is, when a novel object is called a “purple toma,” adults infer that tomas are less commonly purple in general. This inference is captured by a model in the Rational Speech Act framework that posits that listeners reason about speakers’ communicative goals. In Chapter 3, I ask: do children themselves infer that mentioned features are atypical? I find preliminary evidence that 5- to 6-year-old children who reliably respond on our typicality measure can make these inferences, and map out possibilities to further test this question with young children. Overall, this dissertation examines how language does not directly reflect the world, but selectively picks out remarkable facets of it, and what this implies for how adults, children, and language models learn.

INTRODUCTION

An utterance can say much more about the world than its literal interpretation might suggest. For instance, if you hear a colleague say “We should hire a female professor,” you might infer something about the speaker’s goals, the makeup of a department, or even the biases of a field—none of which is literally stated. These inferences depend on recognition that a speaker’s intended meaning can differ from the literal meaning of their utterance, and the process of deriving this intended meaning is called *pragmatics*. Frameworks for understanding pragmatic inference posit that speakers tend to follow general principles of conversation—for instance, that they tend to be relevant, brief, and otherwise helpfully informative (Grice 1975; Sperber and Wilson 1986; Clark 1990). When a speaker deviates from these principles, a listener can reason about the alternative utterances the speaker might have said and infer some intended meaning that goes beyond the literal meaning of their utterance.

Beyond enriching the interpretation of utterances whose literal meaning is known, pragmatic inference is a potentially powerful mechanism for learning about new words and concepts. People can learn the meanings of words by tracking associations between word use and present objects alone (Yu and Smith 2007), but reasoning about a speaker’s intended meaning—not just relating the words they say to objects in the environment—may support more rapid and more accurate learning (Frank, Goodman, and Tenenbaum 2009). For example, Akhtar, Carpenter, and Tomasello (1996) showed that young children can infer the meaning of a new word by using the principle that people tend to remark on things that are new and interesting to them. In this study, an experimenter leaves the room and a new toy emerges in her absence; once she comes back, the toy is familiar to the child but not to the experimenter. When she uses a novel name, “gazzzer,” the child can infer that the word refers to the toy that is novel to the experimenter, and not to other toys the experimenter had already seen. Experiments with adults show that they too can use general principles of informativeness to infer a novel referent’s name (Frank and Goodman 2014).

One potential pragmatic tool for learning about referents is contrastive inference from description. To the extent that communicators strive to be minimal and informative, description should discriminate between the referent and some relevant contrasting set. This contrastive inference is fairly obvious from some types of description, such as some postnominal modifiers: “The door with the lock” clearly implies a contrasting door without one (Ni 1996). The degree of contrast implied by more common descriptive forms, such as prenominal adjectives in English, is less clear: speakers do not always use prenominal adjectives minimally, often describing more than is needed to establish reference (Engelhardt, Baris Demiral, and Ferreira 2011; Mangold and Pobel 1988; Pechmann 1989). Nevertheless, Sedivy et al. (1999) showed that people can use these inferences to resolve referential ambiguity in familiar contexts. When asked to “Pick up the tall cup,” people directed their attention more quickly to the target when a short cup was present, and did so in the period before they heard the word “cup.” Because the speaker would not have needed to specify “tall” unless it was informative, listeners were able to use the adjective to direct their attention to a tall object with a shorter counterpart. Subsequent work using similar tasks has corroborated that people can use contrastive inferences to direct their attention among familiar referents (Sedivy 2003; Aparicio, Xiang, and Kennedy 2016; Ryskin, Kurumada, and Brown-Schmidt 2019).

But what if you didn’t know the meaning of the key words in someone’s utterance—could you use the same kind of contrastive inferences to learn about new words and categories? Suppose a friend asks you to “Pass the tall dax.” Intuitively, your friend must have said the word “tall” for a reason. One possibility is that your friend wants to distinguish the dax they want from another dax they do not. In this case, you might look around the room for two similar things that vary in height, and hand the taller one to them. If, alternatively, you only see one object around whose name you don’t know, you might draw a different inference: this dax might be a particularly tall dax. In this case, you might think your friend

used the word “tall” for a different reason—not to distinguish the dax they want from other daxes around you, but to distinguish the dax they want from other daxes in the world. This would be consistent with data from production studies, in which people tend to describe atypical features more than they describe typical ones (Mitchell, Reiter, and Deemter 2013; Westerbeek, Koolen, and Maes 2015; Rubio-Fernández 2016). For instance, people almost always say “blue banana” to refer to a blue banana, but almost never say “yellow banana” to refer to a yellow one. In each of these cases—when distinguishing the dax from other referents nearby, or from daxes in general—you would have used a pragmatic inference to learn something new about the category of daxes.

This dissertation will explore the ways in which people can learn about new words and categories from contrastive inference, with an eye toward understanding how contrastive inference could help children learn about language and the world it describes. To set the stage for understanding how listeners use contrastive inference, we first need to establish that speakers use adjectives in informative ways.

In Chapter 1, we investigate whether people tend to use adjectives to remark on the atypical features (e.g., “the purple carrot”) rather than the typical features (e.g., “the [orange] carrot”) of things. In a corpus study of caregivers’ speech, we show that caregivers tend to mention atypical rather than typical features of things when speaking to their children. We also show that adults speaking to other adults in naturalistic contexts tend to remark on atypical features rather than typical ones, extending findings from reference game tasks in the lab (Mitchell, Reiter, and Deemter 2013; Westerbeek, Koolen, and Maes 2015; Rubio-Fernández 2016). Finally, we show that children’s own speech mentions atypical more than typical features, and discuss the implications of this finding for our understanding of children’s pragmatic competence.

Given that speech emphasizes atypical features, learning about typicality from language may not be straightforward. In an analysis using language models, we examine whether

it is possible to learn about the typical features of things using the statistical patterns within language alone. To do this, we examine whether two language models that use word co-occurrence to represent word meaning, word2vec and BERT, represent nouns as more similar to or likely to follow their typical adjectives than their atypical adjectives. We find that they have fairly low accuracy on this task: likely because they use associative methods to represent word meaning and their input tends to highlight atypical features, these models represent the relationship between nouns and adjectives poorly. We discuss implications for children’s word learning as well as for language modeling.

In Chapter 2, we establish that adults can use contrastive inferences to learn about a new category’s feature distribution. People use adjectives for multiple communicative purposes: in some cases, an adjective is needed to pick out one object among others in the immediate environment (e.g., “the tall cup” contrasts with a nearby shorter cup, but is not especially tall); in others, it marks atypicality (e.g., “the tall cup” is taller than most cups in general). In this chapter, we use two experiments with adults to show that people can use contrastive inferences to learn about a new category’s feature distribution. People observe instances of novel categories and hear them described (e.g., “Pass me the [green] toma”), and then judge the prevalence of the relevant feature (e.g., how common it is for tomas to be green). People infer that mentioned features are less prevalent than unmentioned ones, and do so even when the feature had to be mentioned to establish reference. We use a model in the Rational Speech Act (RSA) framework to capture people’s judgments, finding that their judgments reflect graded consideration of both reference and conveying typicality as purposes of using an adjective.

In Chapter 3, we present a preliminary study of children’s own contrastive inferences. We test whether children infer that, for example, mentioning that a certain object is tall, blue or spotted implies that other group members are less likely to have those features. However, testing children in this kind of task presents a key difficulty: young children often struggle

with the kinds of scales we use to ask adults about typicality. Our study therefore has two goals: both to examine whether 5- to 6-year-old children can sensibly report typicality on a scale from *few* to *almost all*, and to gather preliminary evidence about their contrastive inferences. We find that though about half of children in this age range struggle with this measure, children who do understand the measure show evidence of contrastive inference. We discuss the implications of this kind of inference for children's learning given the descriptions they hear from caregivers, and the potential unintended consequences of remarking on individuals' traits for children's learning.

CHAPTER 1

PEOPLE TALK MORE ABOUT ATYPICAL THAN TYPICAL FEATURES OF THINGS

Children learn a tremendous amount about the structure of the world around them in just a few short years, from the rules that govern the movement of physical objects to the hierarchical structure of natural categories and even the relational structures among social and cultural groups (Baillargeon 1994; Rogers and McClelland 2004; Legare and Harris 2016). Where does the information driving this rapid acquisition come from? Undoubtedly, a sizeable portion comes from direct experience observing and interacting with the world (Sloutsky and Fisher 2004; Stahl and Feigenson 2015). But another important source of information comes from the language people use to talk about the world (Landauer and Dumais 1997; Rhodes, Leslie, and Tworek 2012). How similar is the information from children’s direct experience to the information available in the language children hear?

Two lines of work suggest that they may be surprisingly similar. One compelling area of work is the comparison of semantic structures learned by congenitally blind children to those of their sighted peers. In several domains that would seem at first blush to rely heavily on visual information, such as verbs of visual perception (e.g., *look*, *see*), blind children and adults make semantic similarity judgments that mirror their sighted peers (Landau, Gleitman, and Landau 2009; Bedny et al. 2019). A second line of evidence supporting the similarity of information in perception and language is the broad success of statistical models trained on language alone in approximating human judgments across a variety of domains (Landauer and Dumais 1997; Mikolov et al. 2013; Devlin et al. 2018). Even more compellingly, models trained on both language and perceptual features for some words can infer the perceptual features of linguistically related words entirely from the covariation of language and perception (Johns and Jones 2012).

Still, there is reason to believe that some semantic features may be harder to learn from

language than these findings suggest. This is because we rarely use language merely to provide running commentary on the world around us; instead, we use language to talk about things that diverge from our expectations or those of our conversational partner (Grice 1975). People tend to avoid being over- or under-informative when they speak. In particular, when referring to objects, people are informative with respect to both the referential context and the typical features of the referent (Westerbeek, Koolen, and Maes 2015; Rubio-Fernández 2016). People tend to refer to an object that is typical of its category with a bare noun (e.g., calling an orange carrot “a carrot”), but often specify when an object has an atypical feature (e.g., “a purple carrot”). Given these communicative pressures, naturalistic language statistics may provide surprisingly little evidence about what is typical (Willits, Sussman, and Amato 2008).

If parents speak to children in this minimally informative way, children may be faced with input that emphasizes atypicality in relation to world knowledge they do not yet have. For things like carrots—which children learn about both from perception and from language—this issue may be resolved by integrating both sources of information. Likely almost all of the carrots children see are orange, and hearing an atypical exemplar noted as a “purple carrot” may make little difference in their inferences about the category of carrots more broadly. But for things to which they lack perceptual access—such as rare objects, unfamiliar social groups, or inaccessible features like the roundness of the Earth—much of what they learn must come from language (Harris and Koenig 2006). If language predominantly notes atypical features rather than typical ones, children may overrepresent atypical features as they learn the way things in the world tend to be.

On the other hand, parents may speak to children differently from the way they speak to other adults. Parents’ speech may reflect typical features of the world more veridically, or even emphasize typical features in order to teach children about the world. Parents alter their speech to children along a number of structural dimensions, using simpler syntax and

more reduplications (Snow 1972). Their use of description may reflect similar alignment to children’s abilities by emphasizing typical feature information children are still learning.

We examine the typicality of adjectives with respect to the nouns they describe in a large, diverse corpus of parent-child interactions recorded in children’s homes to ask whether parents talking to their children tend to use adjectives to mark atypical features. We find that they do: Parents overwhelmingly choose to mention atypical rather than typical features. We also find that parents use adjectives differently over the course of children’s development, noting highly typical features more often to younger children. We additionally compare parents’ speech to a corpus of adult-adult speech and find that parents’ use of description when talking to children is quite similar to adults’ use of description when talking to other adults, and becomes more so as children get older.

We then ask whether the co-occurrence structure of language nonetheless captures typicality information by testing whether language models trained on child-directed speech and adult-directed text capture adjective-noun typicality. We find that relatively little typical feature information is represented in these semantic spaces.

Children’s *own* speech offers a window into how children treat adjectives: do children choose to remark on atypical features themselves? We examine children’s speech in the same corpus of parent-child interactions and find that children too mostly remark on the atypical rather than typical features of things. Though this observational finding cannot provide definitive evidence that children use description to be selectively informative about atypical features, it suggests that even early in life their speech is shaped by adults’ pattern of selective description.

1.1 Adjective typicality

In order to determine whether parents use adjectives mostly to mark atypical features of categories, we analyzed caregiver speech from a large corpus of parent-child interactions, as

well as adult-adult speech as a comparison. We extracted adjectives and the nouns they modified from caregiver speech, and asked a sample of Amazon Mechanical Turkers to judge how typical the property described by each adjective was for the noun it modified. We then examined both the broad features of this typicality distribution and the way it changes over development.

1.1.1 *Corpora*

We used data from the Language Development Project, a large-scale, longitudinal corpus of parent-child interactions recorded in children’s homes. Families were recruited to be representative of the Chicagoland area in both socio-economic and racial composition; all families spoke English at home (Goldin-Meadow et al. 2014). Recordings were taken in the home every 4 months from when the child was 14 months old until they were 58 months old, resulting in 12 timepoints. Each recording was of a 90-minute session in which parents and children were free to behave and interact as they liked.

Our sample consisted of 64 typically-developing children and their caregivers with data from at least 4 timepoints (*mean* = 11.3 timepoints). Together, this resulted in a total of 641,402 parent utterances and 368,348 child utterances.

As an adult-adult speech comparison, we used data from the Conversation Analytic British National Corpus, a corpus of naturalistic, informal conversations in people’s everyday lives (Albert, Ruiter, and Ruiter 2015; Coleman et al. 2012). We excluded any conversations with child participants, for a total of 99,305 adult-adult utterances.

1.1.2 *Stimulus Selection*

We parsed each utterance in our corpora using UDPipe, an automated dependency parser, and extracted adjectives and the nouns they modified. This set contained a number abstract or evaluative adjective-noun pairs whose typicality would be difficult to classify (e.g., “good”–

utterance	pair	rating 1	rating 2	rating 3	mean
especially with wooden shoes.	wooden-shoe	2	2	2	2.00
you like red onions?	red-onion	5	3	4	3.60
the garbage is dirty.	dirty-garbage	7	6	6	6.00

Table 1.1: Sample typicality ratings from three human coders for three adjective-noun pairs drawn from the corpus. Note that means may be slightly different from the mean of the three ratings shown here because some pairs have more than three ratings.

“job”; “little”–“bit”). To resolve this issue, we used human judgments of words’ concreteness to identify and exclude non-concrete adjectives and nouns (Brysbaert, Warriner, and Kuperman 2014). We retained for analysis only pairs in which both the adjective and noun were in the top 25% of concreteness ratings (e.g., “dirty” – “dish”; “green” – “fish”). Additionally, one common adjective that is used abstractly and evaluatively in British English but is concrete in American English (*bloody*) was excluded from the set of pairs from the CABNC.

Our final sample included 6,370 unique adjective-noun pairs drawn from 7,471 parent utterances, 2,775 child utterances, and 1,867 adult-adult utterances. The pairs were combinations of 1,498 distinct concrete nouns and 1,388 distinct concrete adjectives. We compiled these pairs and collected human judgments on Amazon Mechanical Turk for each pair, as described below. Table 1.1 contains example utterances from the final set and typicality judgments from our human raters.

1.1.3 Participants

Each participant rated 35 adjective-noun pairs, and we aimed for each pair to be rated five times, for a total of 910 rating tasks. Participants were allowed to rate more than one set of pairs and were paid \$0.80 per task. Distribution of pairs was balanced using a MongoDB database that tracked how often sets of pairs had been rated. If a participant allowed their task to expire with the task partially complete, we included those ratings and re-recruited the task. Overall, participants completed 32,461 ratings. After exclusions using an attention check that asked participants to simply choose a specific number on the scale, we retained

32,293 judgments, with each adjective–noun pair retaining at least two judgments.

1.1.4 Design and Procedure

To evaluate the typicality of the adjective–noun pairs that appeared in parents’ speech, we asked participants on Amazon Mechanical Turk to rate each pair. Participants were presented with a question of the form “How common is it for a cow to be a brown cow?” and asked to provide a rating on a seven-point scale: (1) never, (2) rarely, (3) sometimes, (4) about half the time, (5) often, (6) almost always, (7) always. We also gave participants the option to select “Doesn’t make sense” if they could not understand what the adjective–noun pair would mean. Pairs that were marked with “Doesn’t make sense” by two or more participants were excluded from the final set of pairs: 1,591 pairs were excluded at this stage, for a final set of 4,779 rated adjective–noun pairs.

1.1.5 Results

We combined the human typicality ratings with usage data from our corpora to examine the extent to which parents, children, and adults speaking to other adults use language to describe typical and atypical features. In our analyses, we token-weighted these judgments, giving higher weight to pairs that occurred more frequently in speech. However, results are qualitatively identical and all significant effects remain significant when examined on a type level.

If caregivers speak informatively to convey what is atypical or surprising in relation to their own sophisticated world knowledge, we should see that caregiver description is dominated by adjectives that are sometimes or rarely true of the noun they modify. If instead child-directed speech privileges redundant information, perhaps to align to young children’s limited world knowledge, caregiver description should yield a distinct distribution dominated by highly typical modifiers. As we predicted, we found that parents’ description




Figure 1.1: Density plots showing parents’ use of atypical and typical adjective-noun pairs across their child’s age.

predominantly focuses on features that are atypical (Figure 1.1).

To confirm this effect statistically, we centered the ratings (i.e. “about half” was coded as 0), and then predicted the rating on each trial with a mixed effect model with only an intercept and a random effect of noun ($\text{typicality} \sim 1 + (1|\text{noun})$). The intercept was reliably negative, indicating that adjectives tend to refer to atypical features of objects ($\beta = -0.85$, $t = -28.611$, $p < .001$). We then re-estimated these models separately for each age in the corpus, and found a reliably negative intercept for every age group (smallest effect $\beta_{14} = -0.684$, $t = -9.063$, $p = < .001$). Even when talking with very young children, caregiver




Figure 1.2: Density plots showing use of atypical and typical adjective-noun pairs by parents speaking to children and adults speaking to other adults.

speech is structured according to the kind of communicative pressures observed in adult-adult conversation in the lab.

To examine whether this holds for naturalistic adult-adult conversation, we performed the same analyses on usage of adjective-noun pairs in adult-adult speech in the Conversation Analytic British National Corpus. The overall distribution of adjective-noun typicality is remarkably similar between child-directed and adult-directed speech (Figure 1.2). Fitting the same mixed-effects model to the adult-directed data, we found that the intercept was reliably negative, indicating that adult-adult speech also predominantly highlights atypical features ($\beta = -0.925$, $t = -29.761$, $p < .001$).

Returning to caregiver speech, while descriptions at every age tended to point out atypical features (as in adult-adult speech), this effect changed in strength over development. As predicted, an age effect added to the previous model was reliably negative, indicating that parents of older children are relatively more likely to focus on atypical features ($\beta = -0.071$,




Figure 1.3: Proportion of caregiver description that is about highly typical features (often, almost always, or always true), as a function of age.

$t = -3.852$, $p = < .001$). In line with the idea that caregivers adapt their speech to their children’s knowledge, it seems that caregivers are more likely to provide description of typical features for their young children, compared with older children. As a second test of this idea, we defined adjectives as highly typical if Turkers judged them to be ‘often’, ‘almost always’, or ‘always’ true. We predicted whether each judgment was highly typical from a mixed-effects logistic regression with a fixed effect of age (log-scaled) and a random effect of noun. Age was a highly reliable predictor ($\beta = -0.688$, $t = -3.78$, $p = < .001$). While children at all ages hear more talk about what is atypically true (Figure 1.1), younger children hear relatively more talk about what is typically true than older children do (Figure 1.3).

What about children themselves—do they tend to remark on the atypical rather than the typical features of things? We analyzed children’s own use of description and found that, following the pattern of parent speech and adult-adult speech, they predominantly mention atypical rather than typical features (Figure 1.4). The fact that children are remarking on atypical features is intriguing, but it would be premature to conclude that they are doing so to be selectively informative. Note also that especially at young ages, children produce few




Figure 1.4: Density plots showing children’s use of atypical and typical adjective-noun pairs across age.

adjective-noun pairs—they are not producing any at 14 months old, our earliest timepoint—so our data on children’s speech is somewhat sparse. We discuss potential interpretations of this finding further in the Conclusion.

1.1.6 Discussion

In sum, we find robust evidence that language is used to discuss atypical, rather than typical, features of the world. Description in caregiver speech seems to largely mirror the usage patterns that we observed in adult-to-adult speech, suggesting that these patterns arise

from general communicative pressures. Interestingly, the descriptions children hear change over development, becoming increasingly focused on atypical features. The higher prevalence of typical descriptors in early development may help young learners learn what is typical; however, even at the earliest point we measured, the bulk of language input describes atypical features.

This usage pattern aligns with the idea that language is used informatively in relation to background knowledge about the world. It may pose a problem, however, for young language learners with still-developing world knowledge. If language does not transparently convey the typical features of objects, and instead (perhaps misleadingly) notes the atypical ones, how might children come to learn what objects are typically like? One possibility is that information about typical features is captured in more complex regularities across many utterances. If this is true, language may still be an important source of information about typicality as children may be able to extract more accurate typicality information by tracking second-order co-occurrence.

1.2 Extracting Typicality from Language Structure

Much information can be gleaned from language that does not seem available at first glance. From language alone, simple distributional learning models can recover enough information to perform comparably to non-native college applicants on the Test of English as a Foreign Language (Landauer and Dumais 1997). Recently, Lewis, Zettersten, and Lupyan (2019) demonstrated that even nuanced feature information may be learnable through distributional semantics alone, without any complex inferential machinery. We take a similar approach to ask whether a distributional semantics model trained on the language children hear can capture typical feature information.

1.2.1 Method

To test this possibility, we trained word2vec—a distributional semantics model—on the same corpus of child-directed speech used in our first set of analyses. Word2vec is a neural network model that learns to predict words from the contexts in which they appear. This leads word2vec to encode words that appear in similar contexts as similar to one another (Firth 1957).

We used the continuous-bag-of-words (CBOW) implementation of word2vec in the `gensim` package (Řehůřek and Sojka 2010). We trained the model using a surrounding context of 5 words on either side of the target word and 100 dimensions (weights in the hidden layer) to represent each word. After training, we extracted the hidden layer representation of each word in the model’s vocabulary—these are the vectors used to represent these words.

If the model captures information about the typical features of objects, we should see that the model’s noun-adjective word pair similarities are correlated with the typicality ratings we elicited from human raters. For a second comparison, we also used an off-the-shelf implementation of word2vec trained on Wikipedia (Mikolov et al. 2018). While the Language Development Project corpus likely underestimates the amount of structure in children’s linguistic input, Wikipedia likely overestimates it.

While word2vec straightforwardly represents what can be learned about word similarity by associating words with similar contexts, it does not represent the cutting edge of language modeling. Perhaps more sophisticated models trained on larger corpora would represent these typicalities better. To test this, we asked how BERT (Devlin et al. 2018) and GPT-3 (Brown et al. 2020) represent typicality. BERT is a masked language model trained on BookCorpus and English Wikipedia, which represents the probability of words occurring in slots in a phrase. We gave BERT phrases of the form “_____ apple”, and asked it the probability of different adjectives filling the empty slot. GPT-3 is a generative language model trained on large quantities of internet text, including Wikipedia, book corpora, and

web page text from crawling the internet. Because it is a generative language model, we can ask GPT-3 the same question we asked human participants directly and it can generate a text response. We prompted the `davinci-text-003` instance of GPT-3 questions of the form: “You are doing a task in which you rate how common it is for certain things to have certain features. You respond out of the following options: Never, Rarely, Sometimes, About half the time, Often, Almost always, or Always. How common is it for a cow to be a brown cow?” Because BERT and GPT-3 have more complex training objectives and are trained on more and different kinds of language than what children hear, results from these models likely do not straightforwardly represent the information available to children in language. However, results from BERT and GPT-3 can indicate the challenges language models face in representing world knowledge when the language people use emphasizes remarkable rather than typical features.

1.2.2 Results

We find that similarities in the model trained on the Language Development Project corpus have near zero correlation with human adjective–noun typicality ratings ($r = 0.05$, $p = .001$). However, our model does capture other meaningful information about the structure of language, such as similarity within part of speech categories. Comparing with pre-existing large-scale human similarity judgements for word pairs, our model shows significant correlations (correlation with `wordsim353` similarities of noun pairs, 0.28; correlation with `simplex` similarities of noun, adjective, and verb pairs, 0.16). This suggests that statistical patterns in child-directed speech are likely insufficient to encode information about the typical features of objects, despite encoding at least some information about word meaning more broadly.

However, the corpus on which we trained this model was small; perhaps our model did not get enough language to draw out the patterns that would reflect the typical features of objects. To test this possibility, we asked whether word vectors trained on a much larger

corpus—English Wikipedia—correlate with typicality ratings. This model’s similarities were significantly correlated with human judgments, although the strength of the correlation was still fairly weak ($r = 0.338$, $p < .001$). How larger and more sophisticated language models fare? Like Wikipedia-trained word2vec, BERT’s probabilities were significantly correlated with human judgments, though weakly so ($r = 0.154$, $p < .001$). However, GPT-3’s ratings were much better aligned with human judgments ($r = 0.574$, $p < .001$).

Similarity judgments produced by our models reflect many dimensions of similarity, but our human judgments reflect only typicality. To account for this fact and control for semantic differences among the nouns in our set, we performed a second analysis in which we considered only the subset of 109 nouns that had both a high-typicality (rated as at least “often”) and a low-typicality (rated as at most “sometimes”) adjective. We then asked whether the word2vec models rated the high-typicality adjective as more similar to the noun it modified than the low-typicality adjective. The LDP model correctly classified 49 out of 109 (0.45%), which was not different from chance ($p = .338$). The Wikipedia-trained word2vec model correctly classified 84 out of 109 (0.771%), which was better than chance according to a binomial test, but not highly accurate ($p = < .001$). Figure 1.5 shows the word2vec models’ similarities for the 109 nouns and their typical and atypical adjectives alongside scaled average human ratings.

The analogous analysis on BERT asks whether the model rates the high-typicality adjective as more likely to come before the noun than the low typicality adjective (e.g., $P(\text{“red”}) > P(\text{“brown”})$ in “_____ apple”). BERT correctly classified 66 out of 109 (0.606%), which is significantly better than chance ($p = .035$). However, BERT’s performance was directionally less accurate than Wikipedia-trained word2vec: though BERT is a more sophisticated model, it does not capture adjective-noun typicality better than word2vec in this analysis. GPT-3 performs much better than BERT and the word2vec models, with 96 out of 109 (0.881%), significantly better than chance by a binomial test ($p = < .001$). Figure 1.6 shows BERT and



Figure 1.5: Plots of word2vec noun-adjective similarities for nouns for which there was at least one atypical adjective (rated at most "sometimes"), and at least one typical adjective (rated at least "often").




Figure 1.6: Plots of BERT and GPT-3 noun-adjective similarities for nouns for which there was at least one atypical adjective (rated at most "sometimes"), and at least one typical adjective (rated at least "often").

GPT-3 ratings for the 109 nouns and their typical and atypical adjectives alongside scaled average human ratings.

1.3 General Discussion

Language provides children a rich source of information about the world. However, this information is not always transparently available: because language is used to comment on the atypical, it does not perfectly mirror the world. Among adult conversational partners whose world knowledge is well-aligned, this principle allows people to converse informatively and avoid redundancy. But between a child and caregiver whose world knowledge is asymmetric, this pressure competes with other demands: what is minimally informative to an adult may be misleading to a child. Our results show that this pressure structures language to create a peculiar learning environment, one in which caregivers predominantly point out the atypical features of things.

How, then, do children learn about the typical features of things? While younger children may gain an important foothold from hearing more description of typical features, they still face language dominated by atypical description. When we looked at more nuanced ways of extracting information from language (which may or may not be available to the developing learner), we found that two word2vec models, one trained on child-directed language and one trained on adult-adult language, did not capture typicality very well. Even BERT, a language model trained on much more text and with a more complex architecture, did not perform better than a Wikipedia-trained word2vec model in reflecting typicality. This may be because these models are designed to capture language statistics, with BERT in particular capturing which words are likely to occur following one another—and as we show in our corpus analyses, adjective-noun pairs that come together often reflect atypicality rather than typicality. Note that a consistent *inverse* relationship—rating high-typicality pairs as *less* similar or *less* probable—would also be evidence that these models capture typicality, but the word2vec models and BERT do not evince this pattern either. However, GPT-3 captured typicality quite well, suggesting that the way people structure language to emphasize atypicality is not necessarily an impediment for much larger models’ representation of typicality. Further

work remains to understand how GPT-3 comes to represent typicality relationships so much better than the smaller models we tested. Overall, a large language model trained on text much greater in quantity and different in quality from child-directed language did capture adjective-noun typicality well, but models with simpler learning mechanisms and language input more similar to what is available to children did not.

Of course, perceptual information from the world may simplify the problem of learning about typicality. In many cases, perceptual information may swamp information from language; children likely see enough orange carrots in the world to outweigh hearing “purple carrot.” It remains unclear, however, how children learn about categories for which they have scarcer evidence. Indeed, language information likely swamps perceptual information for many other categories, such as abstract concepts or those that cannot be learned about by direct experience. If such concepts pattern similarly to the concrete objects analyzed here, children are in a particularly difficult bind.

It is also possible that other cues from language and interaction provide young learners with clues to what is typical or atypical, and these cues are uncaptured by our measure of usage statistics. Caregivers may highlight when a feature is typical by using certain syntactic constructions, such as generics (e.g., “tomatoes are red”). Caregivers may also mark the atypicality of a feature, for example demonstrating surprise. Such cues from language and the interaction may provide key information in some cases; however, given the sheer frequency of atypical descriptors, it seems unlikely that they are consistently well-marked.

Another possibility is that children expect language to be used informatively at a young age. Under this hypothesis, their language environment is not misleading at all, even without additional cues from caregivers. Children as young as two years old tend to use words to comment on what is new rather than what is known or assumed (Baker and Greenfield 1988). Children may therefore expect adjectives to comment on surprising features of objects. If young children expect adjectives to mark atypical features (Horowitz and Frank 2016),

they can use description and the lack thereof to learn more about the world. Our finding that children themselves mostly remark on atypical rather than typical features of things is consistent with this possibility, though does not provide strong evidence that children understand to use description informatively. We will further investigate this question by studying children’s interpretation of adjectives in Chapter 3.

Across our analyses, language is used with remarkable consistency: people talk about the atypical. Though parents might reasonably be broadly over-informative in order to teach their children about the world, this is not the case. This presents a potential puzzle for young learners who have limited world knowledge and limited pragmatic inferential abilities. Perceptual information and nascent pragmatic abilities may help fill in the gaps, but much remains to be explored to link these explanations to actual learning. Communication pressures are pervasive forces structuring the language children hear, and future work must disentangle whether children capitalize on them or are misled by them in learning about the world.

CHAPTER 2

HOW ADULTS USE CONTRASTIVE INFERENCE TO LEARN ABOUT NEW CATEGORIES

When referring to a *big red dog* or a *hot-air balloon*, we often take care to describe them—even when there are no other dogs or balloons around. Speakers use more description when referring to objects with atypical features (e.g., a yellow tomato) than typical ones (e.g., a red tomato; see Chapter 1 and Bergey, Morris, and Yurovsky 2020; Mitchell, Reiter, and Deemter 2013; Westerbeek, Koolen, and Maes 2015; Rubio-Fernández 2016). This selective marking of atypical objects potentially supplies useful information to listeners: they have the opportunity to not only learn about the object at hand, but also about its broader category.

Horowitz and Frank (2016) demonstrated that, combined with other contrastive cues (e.g., “Wow, this one is a zib. This one is a TALL zib”), prenominal adjectives prompted adults and children to infer that the described referent was less typical than one that differed on the mentioned feature (e.g., a shorter zib). This work provided a useful demonstration that adjective use can contribute to inferences about feature typicality, though it did not isolate the effect of adjectives specifically. Their experiments used several contrastive cues, such as prosody (contrastive stress on the adjective: “TALL zib”), demonstrative phrases that may have marked the object as unique (“this one”) and expressions of surprise at the object (“wow”), and participants may have inferred the object was atypical primarily from these cues and not from the adjective. In Chapter 2, we test whether adjective use alone prompts an inference of atypicality with respect to the category’s feature distribution: when you hear “purple toma,” do you infer that *fewer* tomas in general are purple?

If listeners do make contrastive inferences about typicality, it may not be as simple as judging that a described referent is atypical. Description can serve many purposes. If a descriptor is needed to distinguish between two present objects, it may not have been used to mark atypicality. For instance, in the context of a bin of heirloom tomatoes, a speaker who

wants a red one in particular might specify that they want a “red tomato” rather than just asking for a “tomato.” In this case, the adjective “red” is being used contrastively with respect to reference, and not to mark atypicality. If reference is the primary motivator of speakers’ word choice, as implicitly assumed in much research (e.g., Pechmann 1989; Engelhardt, Barış Demiral, and Ferreira 2011; Arts et al. 2011), then people should draw no further inferences once the need for referential disambiguation explains away a descriptor like “red.” On this reference-first view, establishing reference has priority in understanding the utterance, and any further inferences are blocked if the utterance is minimally informative with respect to reference. If, on the other hand, pragmatic reasoning weighs multiple goals simultaneously—here, reference and conveying typicality—people may integrate typicality as just one factor the speaker considers in using description, leading to graded inferences about the referent’s identity and about its category’s features.

In two experiments, we used an artificial language task to set up just this kind of learning situation. We manipulated the contexts in which listeners hear adjectives modifying novel names of novel referents. These contexts varied in how useful the adjective was to identify the referent: some contexts the adjectives were necessary for reference, and in others they were unhelpful. On a *reference-first view*, use of an adjective that was necessary for reference can be explained away and should not prompt further inferences about typicality—an atypicality inference would be blocked. If, on the other hand, people take into account speakers’ multiple reasons for using adjectives without giving priority to reference (the *probabilistic weighing view*), they may alter their inferences about typicality across these contexts in a graded way: if an adjective was necessary for reference, it may prompt slightly weaker inferences of atypicality; if an adjective was redundant with respect to reference, it may be inferred to mark atypicality more strongly. Further, these contexts may also prompt distinct inferences when no adjective is used: for instance, when an adjective is necessary to identify the referent but elided, people may infer that the elided feature is particularly typical. To account for the




Figure 2.1: Experiment 1 stimuli. In the above example, the critical feature is size and the object context is a within-category contrast: the alien on the right has two same-shaped objects that differ in size.

multiple ways context effects might emerge, we analyze both of these possibilities. Overall, we asked whether listeners infer that these adjectives identify atypical features of the named objects, and whether the strength of this inference depends on the referential ambiguity of the context in which adjectives are used.

2.1 Experiment 1

2.1.1 Method

Participants.

240 participants were recruited from Amazon Mechanical Turk. Half of the participants were assigned to a condition in which the critical feature was color (red, blue, purple, or green), and the other half of participants were assigned to a condition in which the critical feature was size (small or big). Participants were paid \$0.30. Participants were told the task was estimated to take 3 minutes and on average took 118 seconds to complete the task (not including reading the consent form).

Stimuli & Procedure.

Stimulus displays showed two alien interlocutors, one on the left side (Alien A) and one on the right side (Alien B) of the screen, each with two novel fruit objects beneath them (Figure 2.1). Alien A, in a speech bubble, asked Alien B for one of its fruits (e.g., “Hey, pass me the big toma”). Alien B replied, “Here you go!” and the referent disappeared from Alien B’s side and reappeared on Alien A’s side. Note that the participants do not make a referent choice in this experiment; the measure of interest is their typicality judgments of the objects’ features, described below.

We manipulated three factors: utterance type, feature type, and context type. We prioritized utterance type as a within-subjects manipulation because it was the central manipulation of interest. We also prioritized context type because another central question was whether context would alter the effect of utterance. We manipulated the critical feature type (color or size) between subjects, to maximize our use of the set of novel stimuli without showing any participant the same novel shape on more than one trial.

Utterance type and context type were fully crossed within subjects. Utterance type had two levels: *adjective noun* (e.g., “Hey, pass me the big toma” or “Hey, pass me the blue toma”) or *noun* (e.g., “Hey, pass me the toma”). Context type had three levels: within-category contrast, between-category contrast, and same feature (Figure 2.2). In the within-category contrast condition, Alien B possessed the target object and another object of the same shape, but with a different feature value (e.g., a big toma and a small toma). In the between-category contrast condition, Alien B possessed the target object and another object of a different shape, and with a different feature value (e.g., a big toma and a small blicket). In the same feature condition, Alien B possessed the target object and another object of a different shape and with the same feature as the target (e.g., a big toma and a big dax). Thus, in the within-category contrast condition, the descriptor was necessary to distinguish the referent; in the between-category contrast condition it was unnecessary but potentially

helpful; and in the same feature condition it was unnecessary and unhelpful.

Note that in all context conditions, the set of objects on screen was the same in terms of the experiment design: there was a target (e.g., big toma), an object with the same shape as the target and a different critical feature (e.g., small toma), an object with a different shape from the target and the same critical feature (e.g., big dax), and an object with a different shape from the target and a different critical feature (e.g., small blicket). Context was manipulated by rearranging these objects such that the relevant referents (the objects under Alien B) differed and the remaining objects were under Alien A. Thus, in each case, participants saw the target object and one other object that shared the target object’s shape but not its critical feature—they observed the same kind of feature distribution of the target object’s category in each trial type.

The particular values of the features were chosen randomly for each trial, and fruits were chosen randomly at each trial from 25 fruit kinds. Ten of the 25 fruit drawings were adapted and redrawn from Kanwisher et al. (1997); we designed the remaining 15 fruit kinds. Each fruit kind had an instance in each of four colors (red, blue, green, or purple) and two sizes (big or small).

Participants completed six trials. After each exchange between the alien interlocutors, they made a judgment about the prevalence of the target’s critical feature in the target object’s category. For instance, after seeing a red blicket being exchanged, participants would be asked, “On this planet, what percentage of blickets do you think are red?” They answered on a sliding scale between zero and 100. In the size condition, participants were asked, “On this planet, what percentage of blickets do you think are the size shown below?” with an image of the target object they just saw available on the screen.

After completing the study, participants were asked to select which of a set of alien words they had seen previously during the study. Four were words they had seen, and four were novel lure words. Participants were dropped from further analysis if they did not respond

to at least 6 of these 8 correctly (above chance performance as indicated by a one-tailed binomial test at the $p = .05$ level). This resulted in excluding 47 participants, leaving 193 for further analysis.

2.1.2 Results

Our key test is whether participants infer that a mentioned feature is less typical than one that is not mentioned. In addition, we tested whether inferences of atypicality are modulated by context. One way to test this is to analyze the interaction between utterance type and context, seeing if the difference between adjective and no adjective utterances is larger when the adjective was highly redundant or smaller when the adjective was necessary for reference.

We analyzed participants' judgments of the prevalence of the target object's critical feature in its category. We began by fitting a maximum mixed-effects linear model with effects of utterance type (adjective or no adjective), context type (within category, between category, or same feature, with between category as the reference level), and critical feature (color or size) as well as all interactions and random slopes of utterance type and context type nested within subject. Random effects were removed until the model converged. The final model included the effects of utterance type, context type, and critical feature and their interactions, and a random slope of utterance type by subject.

This model revealed a significant effect of utterance type ($\beta_{adjective} = -10.22$, $t = -3.374$, $p = .001$), such that prevalence judgments were lower when an adjective was used than when it was not. Participants' inferences did not significantly differ between color and size adjective conditions ($\beta_{size} = 4.728$, $t = 1.455$, $p = .146$). Participants' inferences did not significantly vary by context type ($\beta_{within} = 3.924$, $t = 1.628$, $p = .104$; $\beta_{same} = -1.485$, $t = -0.618$, $p = .537$). There was not a significant interaction between context and presence of an adjective in the utterance ($\beta_{within*adjective} = -1.578$, $t = -0.463$, $p = .644$; $\beta_{same*adjective} = 2.131$, $t = 0.625$, $p = .532$). That is, participants did not significantly adjust their inferences

based on object context, nor did they make differential inferences based on the combination of context and adjective use. However, they robustly inferred that mentioned features were less prevalent in the target’s category than unmentioned features.

This lack of a context effect may be because people do not take context into account, or because they make distinct inferences when an adjective is *not* used: for instance, when an adjective is necessary for reference but elided, people may infer that the unmentioned feature is very typical. This inference would lead to a difference between the *noun* and *adjective noun* utterances in the within-category context, but not because people are failing to attribute the adjective to reference. To account for this possibility, we separately tested whether there are effects of context among just the *noun* trials and just the *adjective noun* trials. In each case, we fit a model with effects of context type and critical feature as well as their interaction and random slopes by subject. Participants did not significantly adjust their inferences by context among only the *noun* trials ($\beta_{within} = 3.945$, $t = 1.469$, $p = .143$; $\beta_{same} = -1.456$, $t = -0.544$, $p = .587$), though numerically they made higher prevalence judgments in the within-category context. That is, we did not find evidence here that people were inferring a feature to be highly typical because it went unmentioned when it was necessary for reference. Participants also did not significantly adjust their inferences by context among only the *adjective noun* trials ($\beta_{within} = 2.434$, $t = 1.159$, $p = .247$; $\beta_{same} = 0.67$, $t = 0.319$, $p = .750$), though their judgments were numerically higher in the within-category context. That is, we did not find evidence that people modulated their typicality inferences based on the referential context among trials where this inference could not have been driven by omission either. Overall, we did not find evidence that participants significantly adjusted their inferences based on context.

2.1.3 Discussion

Description is often used not to distinguish among present objects, but to pick out an object’s feature as atypical of its category. In Experiment 1, we asked whether people would infer that a described feature is atypical of a novel category after hearing it mentioned in an exchange. We found that people robustly inferred that a mentioned feature was atypical of its category, across both size and color description. Further, participants did not use object context to substantially explain away description. That is, even when description was necessary to distinguish among present objects (e.g., there were two same-shaped objects that differed only in the mentioned feature), participants still inferred that the feature was atypical of its category. This suggests that, in the case of hearing someone ask for a “red tomato” from a bin of many-colored heirloom tomatoes, a tomato-naive person would infer that tomatoes are relatively unlikely to be red.

Another interpretation of people’s inferences in the size condition is that they are due to size adjectives being relative gradable adjectives. That is, the phrases “big toma” and “small toma” may inherently carry the meaning “big for a toma” and “small for a toma” (which can be interpreted as an aspect of the adjective’s semantics, not pragmatics; Kennedy 2007; Xiang et al. 2022; Tessler et al. 2020). It is possible to attribute people’s atypicality inferences in the size condition to the relative gradable nature of size adjectives. However, people also made these inferences about color adjectives, which are not relative gradable adjectives. A purely semantic account also might predict that people’s inferences about color and size would be different—for instance, that people would make larger atypicality inferences about size than color—which we do not find. Though the semantics of size adjectives may contribute to people’s inferences of atypicality in the size condition, we find it parsimonious here to explain the color and size inferences by the same mechanism—pragmatic reasoning.

2.1.4 Model

To formalize the inference that participants were asked to make, we developed a model in the Rational Speech Act Framework (RSA, Frank and Goodman 2012). In this framework, pragmatic listeners (L) are modeled as drawing inferences about speakers' (S) communicative intentions in talking to a hypothetical literal listener (L_0). This literal listener makes no pragmatic inferences at all, evaluating the literal truth of a statement (e.g., it is true that a red toma can be called "toma" and "red toma" but not "blue toma"), and chooses randomly among all referents consistent with that statement. In planning their referring expressions, speakers choose utterances that are successful at accomplishing two goals: (1) making the listener as likely as possible to select the correct object, and (2) minimizing their communicative cost (i.e., producing as few words as possible). Note that though determiners are not given in the model's utterances, the assumption that the utterance refers to a specific reference is built into the model structure, consistent with the definite determiners used in the task. Pragmatic listeners use Bayes' rule to invert the speaker's utility function, essentially inferring what the speaker's intention was likely to be given the utterance they produced.

$$Literal : P_{Lit} = \delta(u, r) P(r)$$

$$Speaker : P_S(u|r) \propto \alpha(P_{Lit}(r|u) - C)$$

$$Listener : P_{Learn}(r|u) \propto P_s(u|r) P(r)$$

To allow the Rational Speech Act Framework to capture inferences about typicality, we modified the Speaker's utility function to have an additional term: the listener's expected processing difficulty. Speakers may be motivated to help listeners to select the correct referent

not just eventually but as quickly as possible. People are both slower and less accurate at identifying atypical members of a category as members of that category (Rosch, Simpson, and Miller 1976; Dale, Kehoe, and Spivey 2007). If speakers account for listeners’ processing difficulties, they should be unlikely to produce bare nouns to refer to low typicality exemplars (e.g. unlikely to call a purple carrot simply “carrot”). This is roughly the kind of inference encoded in a continuous semantics Rational Speech Act model (Degen et al. 2020).

We model the speaker as reasoning about the listener’s label verification process. Because the speed of verification scales with the typicality of a referent, a natural way of modeling it is as a process of searching for that particular referent in the set of all exemplars of the named category, or alternatively of sampling that particular referent from the set of all exemplars in that category, $P(r|Cat)$. On this account, speakers want to provide a modifying adjective for atypical referents because the probability of sampling them from their category is low, but the probability of sampling them from the modified category is much higher (a generalization of the size principle, Xu and Tenenbaum 2007). Typicality is just one term in the speaker’s utility, and thus is directly weighed with the literal listener’s judgment and against cost.

If speakers use this utility function, a listener who does not know the feature distribution for a category can use a speaker’s utterance to infer it. Intuitively, a speaker should prefer not to modify nouns with adjectives because they incur a cost for producing an extra word. If they did use an adjective, it must be because they thought the learner would have a difficult time finding the referent from a bare noun alone because of typicality, competing referents, or both. To infer the true prevalence of the target feature in the category, learners combine the speaker’s utterance with their prior beliefs about the feature distribution.

We model the learner’s prior about the prevalence of features in any category as a Beta distribution with two parameters α and β that encode the number of hypothesized prior psuedo-exemplars with the feature and without feature that the learner has previously observed (e.g., one red dax and one blue dax). We assume that the learner believes they have

previously observed one hypothetical psuedo-examplar of each type, which is a weak symmetric prior indicating that the learner expects the target feature value to occur in half of all members of a category on average, but would find many levels of prevalence unsurprising. To model the learner’s direct experience with the category, we add the observed instances in the experiment to these hypothesized prior instances. After observing one member of the category with the target feature value and one without, the listener’s prior is thus updated to be Beta(2, 2).

We used Bayesian data analysis to estimate the posterior mean rationality parameter that participants are using to draw inferences about speakers in both the color and size conditions. The absolute values of these parameters are driven largely by the number of pseudo-exemplars assumed by the listener prior to exposure; however, differences between color and size within the model are interpretable. We found that listeners inferred speakers to be directionally more rational when using size adjectives (0.887 [0.626, 1.134]) than color adjectives (0.604 [0.367, 0.833]), but the two inferred confidence intervals were overlapping, suggesting that people treated size and color adjectives similarly when making inferences about typicality.

Figure 2.2 shows the predictions of our Rational Speech Act model compared to empirical data from participants. The model captures the trends in the data correctly, inferring that the critical feature was less prevalent in the category when it was mentioned (e.g., “red dax”) than when it was not mentioned (e.g., “dax”). The model also infers the prevalence of the critical feature to be numerically higher in the within-category condition, like people do. That is, in the within-category condition when an adjective is used to distinguish between referents, the model thinks that the target color is slightly less atypical. When an adjective would be useful to distinguish between two objects of the same shape but one is not used, the model infers that the color of the target object is slightly more typical.

Overall, our model captures the inference people make: when the speaker mentions a

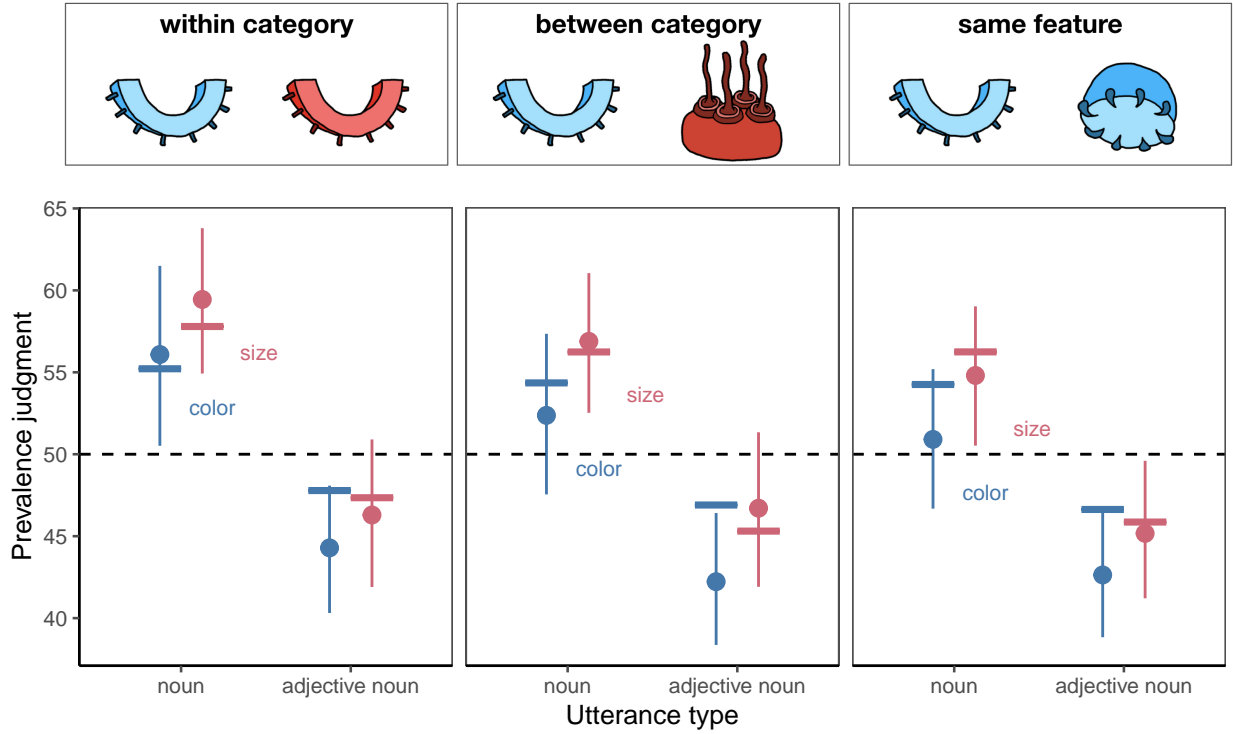


Figure 2.2: Participants’ prevalence judgments from Experiment 1, along with our model predictions. Participants consistently judged the target object as less typical of its category when the referent was described with an adjective (e.g., "Pass me the blue toma") than when it was not (e.g., "Pass me the toma"). This inference was not significantly modulated by object context (examples shown above each figure panel). Points indicate empirical means; error bars indicate 95% confidence intervals computed by non-parametric bootstrapping. Solid horizontal lines indicate model predictions.

feature (e.g., “the blue dax”), people infer that the feature is less typical of the category (daxes are less likely to be blue in general). It further captures that when the object context requires an adjective for successful reference, people weaken this atypicality inference only slightly, if at all. In contrast to a reference-first view, which predicts that these two kinds of inferences would trade off strongly—that is, using an adjective that is necessary for reference would block the inference that it is marking atypicality—the model captures the graded way in which people consider these two communicative goals.

2.2 Experiment 2

In Experiment 1, we established that people can use contrastive inferences to make inferences about the feature distribution of a novel category. Additionally, we found that these two inferences do not seem to trade off substantially: even if an adjective is necessary to establish reference, people infer that it also marks atypicality. To strengthen our findings in a way that would allow us to better detect potential trade-offs between these two types of inference, in Experiment 2 we conducted a pre-registered replication of Experiment 1 with a larger sample of participants. In addition, we tested how people’s prevalence judgments from utterances with and without an adjective compare to their null inference about feature prevalence by adding a control utterance condition: an alien utterance, which the participants could not understand. This also tests the model assumption we made in Experiment 1: that after seeing two exemplars of the target object with two values of the feature (e.g., one green and one blue), people’s prevalence judgments would be around 50%. In addition to validating this model assumption, we more strongly tested the model here by comparing predictions from same model, with parameters inferred from the Experiment 1 data, to data from Experiment 2. Our pre-registration of the method, recruitment plan, exclusion criteria, and analyses can be found on the Open Science Framework: <https://osf.io/s8gre>.

2.2.1 Method

Participants.

A pre-registered sample of 400 participants was recruited from Amazon Mechanical Turk. Half of the participants were assigned to a condition in which the critical feature was color (red, blue, purple, or green), and half of the participants were assigned to a condition in which the critical feature was size (small or big).

Stimuli & Procedure.

The stimuli and procedure were identical to those of Experiment 2, with the following modifications. Two factors, utterance type and object context, were fully crossed within subjects. Object context had two levels: within-category contrast and between-category contrast. In the within-category context condition, Alien B possessed the target object and another object of the same shape, but with a different value of the critical feature (color or size). In the between-category contrast condition, Alien B possessed the target object and another object of a different shape, and with a different value of the critical feature. Thus, in the within-category contrast condition, an adjective is necessary to distinguish the referent; in the between-category contrast condition it is unnecessary but potentially helpful. There were three utterance types: adjective, no adjective, and alien utterance. In the two alien utterance trials, the aliens spoke using completely unfamiliar utterances (e.g., “Zem, noba bi yix blicket”). Participants were told in the task instructions that sometimes the aliens would talk in a completely alien language, and sometimes their language will be partly translated into English. To keep participants from making inferences about the content of the alien utterances using the utterance content of other trials, both alien language trials were first; other than this constraint, trial order was random. We manipulated the critical feature type (color or size) between subjects.

After completing the study, participants were asked to select which of a set of alien words they had seen previously during the study. Four were words they had seen, and four were novel lure words. Participants were dropped from further analysis if they did not meet our pre-registered criteria of responding to at least 6 of these 8 correctly (above chance performance as indicated by a one-tailed binomial test at the $p = .05$ level) and answering all four color perception check questions correctly. Additionally, six participants were excluded because their trial conditions were not balanced due to an error in the run of the experiment. This resulted in excluding 203 participants, leaving 197 for further analysis.

In our pre-registration, we noted that we anticipated high exclusion rates, estimating that approximately 150 people per condition would be sufficient to test our hypotheses.

2.2.2 Results

We began by fitting a pre-registered maximum mixed-effects linear model with effects of utterance type (alien utterance, adjective, or no adjective; alien utterance as reference level), context type (within category or between category), and critical feature (color or size) as well as all interactions and random slopes of utterance type and context type nested within subject. Random effects were removed until the model converged, which resulted in a model with all fixed effects, all interactions and a random slope of utterance type by subject. The final model revealed a significant effect of the no adjective utterance type compared to the alien utterance type ($\beta = 7.476$, $t = 2.798$, $p = .005$) and no significant effect of the adjective utterance type compared to the alien utterance type ($\beta = -0.641$, $t = -0.244$, $p = .808$). The effects of context type (within-category or between-category) and adjective type (color or size) were not significant ($\beta_{within} = -2.699$, $t_{within} = -1.228$, $p_{within} = .220$; $\beta_{size} = 4.435$, $t_{size} = 1.33$, $p_{size} = .185$). There were marginal interactions between the adjective utterance type and the size condition ($\beta = -6.561$, $t = -1.724$, $p = .086$), the adjective utterance type and the within-category context ($\beta = 5.767$, $t = 1.856$, $p = .064$), and the no adjective utterance type and the within-category context ($\beta = 5.573$, $t = 1.793$, $p = .073$). No other effects were significant or marginally significant. Thus, participants inferred that an object referred to in an intelligible utterance with no description was more typical of its category on the target feature than an object referred to with an alien utterance. Participants did not substantially adjust their inferences based on the object context. The marginal interactions between the within-category context and both the adjective and no adjective utterance types suggest that people might have judged the target feature as slightly more prevalent in the within-category context when intelligible utterances (with a bare noun

or with an adjective) were used compared to the alien utterance. If people are discounting their atypicality inferences when the adjective is necessary for reference, we should expect them to have slightly higher typicality judgments in the within-category context when an adjective is used, and this marginal interaction suggests that this may be the case. However, since typicality judgments in the no adjective utterance type are also marginally greater in the within-category context, and because judgments in the alien utterance conditions (the reference category) also directionally move between the two context conditions, it is hard to interpret whether this interaction supports the idea that people are discounting their typicality judgments based on context.

Given that interpretation of these results with respect to the alien utterance condition can be difficult, we pre-registered a version of the same full model excluding alien utterance trials with the no adjective utterance type as the reference level. This model revealed a significant effect of utterance type: participants' prevalence judgments were lower when an adjective was used than when it was not ($\beta = -8.117$, $t = -3.463$, $p = .001$). No other effects were significant. This replicates the main effect of interest in Experiment 1: when an adjective is used in referring to the object, participants infer that the described feature is less typical of that object's category than when the feature goes unmentioned. It also shows that the possibility that people may discount their typicality judgments based on context (suggested by the marginal interaction described above) is not supported when we compare the adjective and no adjective utterance types directly. In the Supplemental Materials, we report two more pre-registered tests of the effect of utterance type alone on prevalence judgments whose results are consistent with the fuller models reported here.

As in Experiment 1, our test of whether participants' inferences are modulated by context is potentially complicated by people making distinct inferences when an adjective is necessary but *not* used. Thus, we additionally tested whether participants' inferences varied by context among only utterances with an adjective by fitting a model with effects of context and

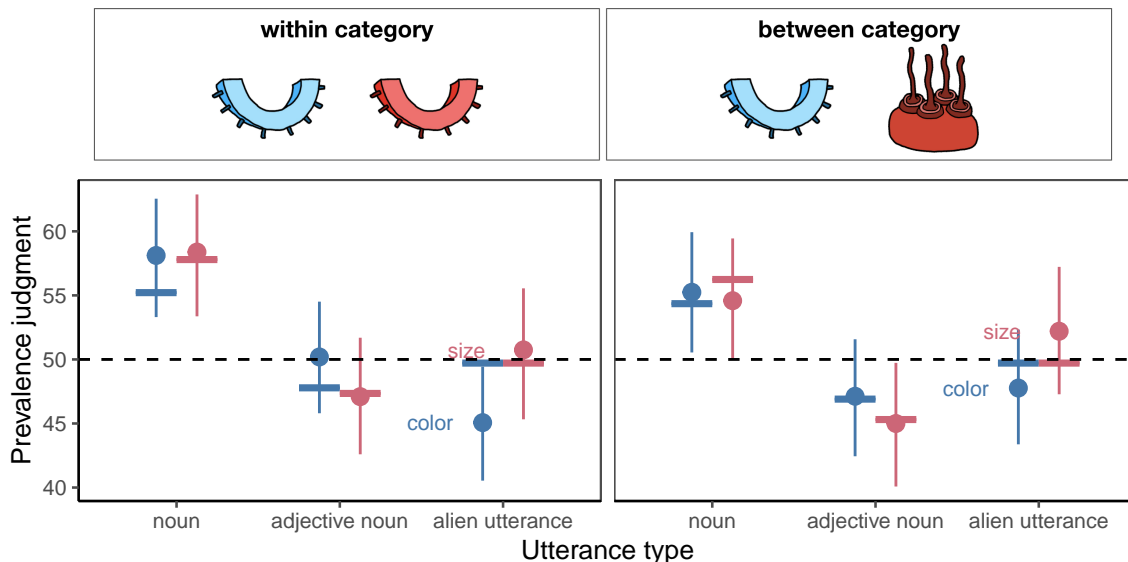


Figure 2.3: Participants’ prevalence judgments in Experiment 2, with model predictions using the parameters estimated in Experiment 1 (horizontal lines).

adjective type and their interaction, as well as random slopes by subject (not pre-registered). Participants’ inferences did not significantly differ by context ($\beta_{within} = 3.068$, $t_{within} = 1.696$, $p_{within} = .091$). Thus, participants’ inferences did not significantly differ between contexts, whether tested by the interaction between utterance type and contexts or by the effect of context among only utterances with an adjective.

2.2.3 Model

To validate the model we developed for Experiment 1, we compared its estimates using the previously fit parameters to the new data from Experiment 2. As shown in Figure 2.3, the model predictions were well aligned with people’s prevalence judgments. In addition, in Experiment 1, we fixed the model’s prior beliefs about the prevalence of the target object’s color or size to be centered at 50% because the model had seen one pseudo-exemplar of the target color/size, and one psuedo-exemplar of the non-target color/size. In Experiment 2, we aimed to estimate this prior empirically in the alien utterance condition, reasoning that people could only use their prior to make a prevalence judgment (as we asked the model

to do). In both the color and size conditions, people’s judgments indeed varied around 50%, although in the color condition they were directionally lower. This small effect may arise from the fact that size varies on a scale with fewer nameable points (e.g., objects can be big, medium-sized or small) whereas color has many nameable alternatives (e.g., red, blue, green, etc.). Thus, the results of Experiment 2 confirm the modeling assumptions we made in estimating people’s prior beliefs, and further validate the model we developed as a good candidate model for how people simultaneously draw inferences about speakers’ intended referents and the typicality of these referents. That is, when people think about why a speaker chose their referring expression, they consider the context of not only present objects, but also the broader category to which the referent belongs.

2.2.4 Discussion

In Experiment 2, we replicated the main finding of interest in Experiment 1: when a novel object’s feature is described, people infer that the feature is rarer of its category than when it goes unmentioned. Again, this effect was consistent across both size and color adjectives, and people did not substantially adjust this inference based on how necessary the description was to distinguish among potential referents. We also added an alien language condition, in which the entire referring expression was unintelligible to participants, to probe people’s priors on feature typicality. We found that in the alien language condition, people judged features to be roughly between the adjective utterance and no adjective utterance conditions, and significantly different from the no adjective utterance condition. In the alien language condition, people’s prevalence judgments were roughly around our model’s prevalence judgments (50%) after observing the objects on each trial and before any inferences about the utterance.

The similarity of people’s prevalence judgments in the alien language condition and the adjective condition raises the question: is this effect driven by an atypicality inference in the

adjective conditions, or a *typicality* inference when the feature is unmentioned? Our results suggest that it is a bit of both. When someone mentions an object without extra description, the listener can infer that its features are likely more typical than their prior; when they use description, they can infer that its features are likely less typical. Because using an extra word—an adjective—is generally not thought of as the default way to refer to something, this effect is still best described as a contrastive inference of *atypicality* when people use description. However, the fact that people infer high typicality when an object is referred to without description suggests that, in some sense, there is no neutral way to refer: people will make broader inferences about a category from even simple mentions of an object.

2.3 Conclusion

In Chapter 1, we established that people tend to mention atypical rather than typical features. In this chapter, we showed that adults make appropriate pragmatic inferences given how speakers describe: they infer that a mentioned feature is likely to be less typical of the mentioned category. However, the ability to learn about new categories using contrastive inference most obviously serves budding language learners—children. To fully appreciate the potential of these inferences to allow people to learn about the world, we must study their development, which we will turn to in Chapter 3.

CHAPTER 3

HOW CHILDREN USE CONTRASTIVE INFERENCE TO LEARN ABOUT NEW CATEGORIES

The speech children hear mentions more atypical than typical features. Depending on children's pragmatic abilities, this input could provide helpful information or pose a misleading challenge as children learn about the world. If children are able to make the contrastive inference that description tends to pick out atypical features, they could use description to go beyond learning about what they directly experience. If, on the other hand, they merely associate the mentioned feature with the mentioned category, they may mistakenly learn that atypical features are more common than they actually are.

In general, children's pragmatic abilities are thought to undergo prolonged development, not reaching adult-like performance until well into schooling age. The most thoroughly studied pragmatic inference in children, scalar implicature, tells a bleak story about children's ability to make pragmatic inferences at a young age. Scalar implicature is the phenomenon in which use of a weak scalar term ('some,' 'might') implies that a stronger scalar term ('all,' 'must') is not true—for example, "I ate some of the cookies" implies I did not eat all of them. This inference can be derived by reasoning that had the speaker meant the stronger meaning, they would have used the stronger term. Adults consistently interpret the word 'some' to mean 'some but not all,' rating the use of 'some' as unnatural when 'all' is applicable and taking longer to respond to such instances (Bott and Noveck 2004; Degen and Tanenhaus 2015). Until at least the age of 5 and in some tasks up to 10 years old, children fail to limit the use of 'some' in this way, accepting 'some' as a descriptor when 'all' is true (Noveck 2001; Papafragou and Musolino 2003). This deficit is found in a range of measures, from acceptability judgments to eye-tracking (Huang and Snedeker 2009). Later work has found that children likely lack this ability because they fail to activate alternative descriptions, so cannot reason that the speaker should have said 'all' and not 'some' if all

is true (Barner, Brooks, and Bale 2011), and because they lack a meta-understanding of these tasks (Papafragou and Musolino 2003). When given supportive context, like named alternatives or training on the task, 4- and 5-year-olds improve at these implicatures (Barner, Brooks, and Bale 2011; Papafragou and Musolino 2003; Foppolo, Guasti, and Chierchia 2012). However, across experiments, performance on scalar implicature remains fragile well into school age.

Contrastive inference from description, however, may be a more accessible form of pragmatic inference because the relevant alternatives are more easily accessible. In the case of using contrastive inference to resolve reference (e.g., “the tall...” prompts looking to a tall object with a shorter counterpart), the relevant alternatives are available in the environment. By the age of 5, children can use contrastive inferences to direct their attention among familiar present objects (Huang and Snedeker 2008), and when given extra time to orient to the referent, show budding abilities by the age of 3 (Davies et al. 2021). Description paired with other contrastive cues can allow children to restrict reference among novel objects or objects with novel properties, though imperfectly (Gelman and Markman 1985; Diesendruck, Hall, and Graham 2006).

What about when the contrasting set is not available in the environment, but is the referent’s category? Preliminary evidence also suggests that contrastive inference about typicality may be possible for young children. When paired with other contrastive cues, 4-year-olds can make inferences about novel object typicality, reasoning that “the TALL zib” suggests other zibs are generally shorter (Horowitz and Frank 2016). This work provided a useful demonstration that adjective use can contribute to inferences about feature typicality, though it did not isolate the effect of adjectives specifically. Their experiments used several contrastive cues, such as prosody (contrastive stress on the adjective: “TALL zib”), demonstrative phrases that may have marked the object as unique (“this one”) and expressions of surprise at the object (“wow”), and participants may have inferred the object was atypical

primarily from these cues and not from the adjective. Further, these experiments used a forced-choice measure that does not allow a precise estimate of how much children’s typicality judgments shift from adjective use. Thus, in this experiment, we set out to develop a task that would isolate the effect of adjective use and measure children’s typicality judgments in a more graded way.

In this chapter, we report an exploratory study of children’s abilities to make contrastive inferences about typicality. To do this, we used a task similar to those done by adults in Chapter 2, having children observe novel categories and make inferences about the typicality of their features. We study 5- to 6-year-old children, an age at which key pragmatic abilities are developing and when children can use contrastive inferences to direct their attention among familiar referents. Because children at this age struggle to explicitly reason about and report proportions (see Boyer, Levine, and Huttenlocher 2008 for a review), we will have children report their typicality judgments with the help of visual depictions of *few*, *some*, *most*, and *almost all* objects having a feature. The purpose of this exploratory study is both to see whether children can make sensible responses on this measure and to gather preliminary evidence about children’s contrastive inferences.

3.1 Method

3.1.1 *Participants.*

We recruited 30 5–6-year-old children raised with 90% or greater English language exposure to participate in this task. Children were recruited from a database with mostly families living in the Chicago area, and some families living elsewhere in the United States, and the study was conducted remotely on Zoom. Data from one participant was excluded due to connection difficulties in the call. In the final sample, 15 5-year-olds and 14 6-year-olds participated.




Figure 3.1: An example of the novel objects shown on a trial. In each trial, two objects of the same shape and differing on the critical feature were shown sequentially. In adjective noun trials, the critical feature was mentioned for the object that had it (e.g., the wide toma was called a “wide toma”) and in noun trials, no features were mentioned (e.g., both tomas were just called a “toma”).




Figure 3.2: An example of the prevalence judgment children were asked to make. Children chose between clouds of novel objects representing few, some, most, and almost all of the novel category having the feature. The experimenter asked, e.g., “Let’s think about all of the blickets on this planet. How many blickets do you think are spotted? Few of the blickets, some of the blickets, most of the blickets, or almost all of the blickets?”

3.1.2 Design and Procedure.

Children participated in a novel object learning task in which they observed novel objects and made inferences about them. They were introduced to an alien named Blip, who would show things from her planet. Blip’s utterances were presented both as recorded audio and displayed in a text bubble on the screen. In each trial, Blip first said “Let’s see what I have...” and then sequentially showed two objects with the same name and shape. The two objects differed on the critical feature. In *adjective noun* trials, the critical feature was mentioned (e.g., one object was labeled “It’s a blicket” and the other was labeled “It’s a striped blicket”); in *noun* trials, the critical feature was not mentioned (e.g., one object was labeled “It’s a blicket” and the other was also labeled “It’s a blicket”).

After each trial, children were asked to make a judgment about the prevalence of the critical feature in the novel category. For instance, they were asked, “Let’s think about all of the blickets on this planet. How many blickets do you think are spotted?” There were four options on the screen, each a cloud of six of the same shape of novel object, with differing proportions having the critical feature and in color (and the remaining objects without the feature and in grey). The options were *Few* (1/6 with feature), *Some* (2/6 with feature), *Most* (4/6 with feature), and *Almost All* (5/6 with feature). After asking the question, the experimenter said the options: “Few of the blickets, some of the blickets, most of the blickets, or almost all of the blickets?” Children responded verbally. If they paused or seemed uncertain, the experimenter repeated the options. If the child preferred to point to the option on the screen (as happened with one participant), the experimenter asked the child’s parent to report the option they pointed to.

There were six trials in total. Half of trials were *adjective noun* trials and half were *noun* trials, and this factor was crossed with the feature type: size (wide or tall), color (blue or red), and pattern (spotted or striped). At each trial, the novel object shape and novel object name were randomly assigned out of a set of six names (modi, blicket, wug, toma, gade, or

sprock) and shapes. The ordering of two objects in each trial (one with the critical feature and one without) was random.

Before the main task, children did two practice trials with familiar objects to establish that they understood the response measure. The two practice questions were: “Let’s think about all of the cookies in the world. How many cookies do you think are square?” and “Let’s think about all of the bananas in the world. How many bananas do you think are yellow?” They responded on the same scale used in the main task trials.

3.2 Performance on practice trials

Children’s performance on the two practice trials with familiar objects can help give us a sense of whether they understand the typicality measure in this task. If the children understand this measure, we expect them to report that bananas are more commonly yellow than cookies are square. Out of 29 participants, 15 (0.517%) rated bananas to be more commonly yellow than cookies are square (0.4% of 5-year-olds and 0.643% of 6-year-olds). That is, many children, especially the 5-year-olds, either did not understand this measure well or did not believe that cookies are not typically square and bananas are typically yellow. Below, we will report the results of the main task both for all children and, separately, for just the children who performed correctly on the familiar practice trials to see whether there is evidence for contrastive inference among children who understood the measure.

3.3 Results

Our key question is whether children make different inferences when an object’s feature is mentioned than when it is not. To test this question, we fit a linear regression with children’s prevalence choices as the outcome (coded as *few* = 1, *some* = 2, *most* = 3, and *almost all* = 4) and utterance type (*noun* vs. *adjective noun*), feature type (color, size, or pattern),




Figure 3.3: Children’s prevalence judgments across utterance conditions and feature types.

and their interaction as predictors, as well as a random intercept by subject. The effect of utterance type was marginally significant: children’s prevalence judgments were marginally lower when there was an adjective in the utterance ($\beta = -0.552$, $t = -1.951$, $p = 0.053$). Effects of feature type were not significant ($\beta_{pattern} = -0.448$, $t = -1.585$, $p = 0.115$; $\beta_{size} = -0.276$, $t = -0.975$, $p = 0.331$), nor were interactions between utterance type and feature type ($\beta_{adjective-noun*pattern} = 0.655$, $t = 1.638$, $p = 0.104$; $\beta_{adjective-noun*size} = 0.483$, $t = 1.207$, $p = 0.23$). Though effects of feature type and the interaction between utterance type and feature type are not significant, visually examining the plotted data, the overall marginal effect of utterance type seems to be driven by the color condition. Overall, we find weak evidence that children infer that mentioned features are less typical. Children’s prevalence judgments are shown in Figure 3.3.

Based on their performance in the practice trials, it seems that many children did not understand the prevalence measure well. We can separately test the performance of chil-




Figure 3.4: Prevalence judgments among only children who answered the practice trials correctly. These children rate features to be less prevalent when they are mentioned with an adjective.

dren who correctly answered the practice trials to see whether children who understand the measure demonstrate contrastive inference. We fit the same model specification to only children who rated bananas to be yellow more typically than cookies are square. Among these children, there is a significant effect of utterance type, such that they infer that mentioned features are less typical ($\beta = -1.133$, $t = -3.069$, $p = .003$). Effects of feature type were not significant ($\beta_{pattern} = -0.467$, $t = -1.264$, $p = 0.21$; $\beta_{size} = -0.2$, $t = -0.542$, $p = 0.59$), nor were interactions between utterance type and feature type ($\beta_{adjective-noun*pattern} = 0.867$, $t = 1.66$, $p = 0.101$; $\beta_{adjective-noun*size} = 0.933$, $t = 1.787$, $p = 0.078$). The utterance effect is also directionally present across all three feature conditions. Children who performed correctly on familiar trials judged mentioned features to be less typical (Figure 3.4).

3.4 Discussion

In this chapter, we ask how children develop the inference that that when a feature of a novel category is mentioned, that feature is likely to be atypical of the category. One possibility is that children simply associate the words, features and categories that are salient in an instance of reference. This would lead children to think a mentioned feature is representative or typical of the mentioned category. Another possibility is that children make the kind of contrastive inference adults make, inferring that the mentioned feature is atypical. In an exploratory study, we found suggestive evidence that 5–6-year-old children are not making an associative inference, and are directionally making a contrastive inference about mentioned features. Further, children who performed correctly on practice trials with familiar objects made significantly lower typicality judgments about mentioned features. However, judging typicality is difficult for young children, and participants struggled with our measure overall. Evidence from this task is only preliminary, and calls for confirmatory tests with larger sample sizes and for the development of measures that are more sensible for young children.

CONCLUSION

This dissertation examines how speakers selectively describe remarkable features and how listeners use this selective description to learn more about the world. In doing so, it inverts the framework that has positioned pragmatic inference as augmenting literal meaning that is already known, instead considering how people can use pragmatics to learn more about the semantics of unfamiliar things.

To understand how people use description to learn about the world, we first must know how description is used. Chapter 1 illustrates how caregivers use description in speaking to children, as well as establishing how adults use description when speaking to other adults and how children themselves use description. We find that parents predominantly mention atypical rather than typical features when speaking to children, as do adults when speaking to other adults.

We also examined how children themselves use description, and found that they mostly talk about the atypical features of things. There are several language-generating processes that may explain children’s use of description. One possibility is that children understand that description is used to draw a distinction between the described thing and some relevant alternatives—that they are using description informatively to highlight atypicality. Another possibility is that children are broadly reflecting the distribution of adjective-noun usage in their parents’ speech, simply by producing the kinds of adjective-noun pairs they have heard before. A third is that their pattern of description is largely explained by local mimicry—that children are directly repeating back adjectives and nouns their parents used recently in conversation. More focused corpus analyses, as well as experiments eliciting children’s adjective production, are necessary to distinguish between these possibilities.

The pattern of description we find in parents’ speech to children is consistent with the idea that people use language informatively with relation to background knowledge of the world, rather than giving veridical running commentary on the world’s features. This finding

raises questions about how children use description to learn, given that so many accounts of language learning rest on children forming associations among co-occurring words, features, and concepts. To test what kind of typicality information is derivable from language alone, we investigated whether language models that use associative learning among words can extract typical feature information from language. We find that simpler distributional semantics models do poorly in distinguishing between the typical and atypical features of nouns, with implications both for associative accounts of children’s language learning and for language modeling. However, a large language model with a more complex architecture and access to more and different language input than children receive was able to capture adjective-noun typicality fairly well. Overall, our findings highlight the complexity of learning about the world from language that describes it selectively.

However, perhaps people—unlike simpler associative language models—know that language is used to selectively remark on the world, and can use this fact to learn about the unfamiliar. In Chapter 2, we investigated how adults make inferences about novel object categories, and found that they can use description to infer that a described feature is atypical. Further, even when description may have been used for another purpose—to establish reference—people make inferences about typicality. We find that a model that considers the utility of utterances with respect to reference and typicality captures people’s inferences. Much prior work has only considered the use of description in distinguishing between present referents (Pechmann 1989; Engelhardt, Barış Demiral, and Ferreira 2011; Mangold and Pobel 1988), and even work that has incorporated typicality has focused on reference as the primary goal of description (Sedivy 2003; Mitchell, Reiter, and Deemter 2013; Westerbeek, Koolen, and Maes 2015; Rubio-Fernández 2016). Our findings emphasize that conveying typicality is likely a central factor in referring, and inferences about typicality are not secondary to or blocked by the purpose of establishing reference. Further, pragmatics is generally conceived of as a layer of meaning that only emerges on top of a more stable semantics; these studies

underscore that when semantic meaning is uncertain, people can use pragmatics to resolve it.

The ability to exploit description to learn more about the world than one has observed directly is most useful to people who are still rapidly learning—children. In Chapter 3, we investigated how 5- to 6-year-old children make contrastive inferences about typicality. The results of our preliminary experiment show that it is difficult to elicit graded typicality judgments from children. However, children who understand our typicality measure do not make associative inferences in this task; rather, we find weak, preliminary evidence that these children make contrastive inferences. Taken together with evidence from our corpus analysis, this preliminary study suggests that by the age of 5 or 6 children are *not* making associative inferences about the atypical adjective-noun pairs they hear, and may be making contrastive inferences instead. However, further work with better measures is necessary to confirm this finding and examine how younger children interpret the description they hear.

The core computation in pragmatic inference is reasoning about alternatives—things the speaker could have said and did not. Given that others are reasoning about these alternatives, no choice is neutral. In the studies in Chapter 2, for instance, using an adjective in referring to an object led people to infer that the feature described by that adjective was less typical than if it had not been mentioned. But, conversely, *not* using an adjective led them to think that the feature was more typical than if they could not understand the meaning of the utterance at all—all communicative choices leak one’s beliefs about the world. This has implications not only for learning about novel concrete objects, as people did here, but for learning about less directly accessible entities such as abstract concepts and social groups. These inferences can be framed positively, as ways for learners to extract additional knowledge that was not directly conveyed, but can also spread beliefs that the speaker does not intend. The principle that people speak informatively is simple, but it holds unintuitive consequences—among speakers and listeners, humans and machines, adults and children—

for describing and learning about the world.

REFERENCES

- Akhtar, Nameera, Malinda Carpenter, and Michael Tomasello. 1996. “The Role of Discourse Novelty in Early Word Learning.” *Child Development* 67 (2): 635–45. <https://doi.org/10.1111/j.1467-8624.1996.tb01756.x>.
- Albert, Saul, Laura E. de Ruiter, and J. P. de Ruiter. 2015. “CABNC: The Jeffersonian Transcription of the Spoken British National Corpus.”
- Aparicio, Helena, Ming Xiang, and Christopher Kennedy. 2016. “Processing Gradable Adjectives in Context: A Visual World Study.” In *Semantics and Linguistic Theory*, 25:413–32.
- Arts, Anja, Alfons Maes, Leo G. M. Noordman, and Carel Jansen. 2011. “Overspecification in Written Instruction.” *Linguistics* 49 (3): 555–74.
- Baillargeon, Renee. 1994. “How Do Infants Learn About the Physical World?” *Current Directions in Psychological Science* 3 (5): 133–40.
- Baker, Nancy D., and Patricia M. Greenfield. 1988. “The Development of New and Old Information in Young Children’s Early Language.” *Language Sciences* 10 (1): 3–34.
- Barner, David, Neon Brooks, and Alan Bale. 2011. “Accessing the Unsaid: The Role of Scalar Alternatives in Children’s Pragmatic Inference.” *Cognition* 118 (1): 84–93.
- Bedny, Marina, Jorie Koster-Hale, Giulia Elli, Lindsay Yazzolino, and Rebecca Saxe. 2019. “There’s More to ‘Sparkle’ Than Meets the Eye: Knowledge of Vision and Light Verbs Among Congenitally Blind and Sighted Individuals.” *Cognition* 189: 105–15.
- Bergey, Claire, Benjamin Morris, and Daniel Yurovsky. 2020. “Children Hear More About What Is Atypical Than What Is Typical.” PsyArXiv. <https://doi.org/10.31234/osf.io/5wvu8>.
- Bott, Lewis, and Ira A. Noveck. 2004. “Some Utterances Are Underinformative: The Onset and Time Course of Scalar Inferences.” *Journal of Memory and Language* 51 (3): 437–57.
- Boyer, Ty W., Susan C. Levine, and Janellen Huttenlocher. 2008. “Development of Proportional Reasoning: Where Young Children Go Wrong.” *Developmental Psychology* 44:

- 1478–90. <https://doi.org/10.1037/a0013110>.
- Brown, Tom B., Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, et al. 2020. “Language Models Are Few-Shot Learners.” arXiv. <https://doi.org/10.48550/arXiv.2005.14165>.
- Brysbaert, Marc, Amy Beth Warriner, and Victor Kuperman. 2014. “Concreteness Ratings for 40 Thousand Generally Known English Word Lemmas.” *Behavior Research Methods* 46 (3): 904–11.
- Clark, Eve V. 1990. “On the Pragmatics of Contrast.” *Journal of Child Language* 17 (2): 417–31. <https://doi.org/10.1017/S0305000900013842>.
- Coleman, John, Ladan Baghai-Ravary, John Pybus, and Sergio Grau. 2012. “Audio BNC: The Audio Edition of the Spoken British National Corpus.”
- Dale, Rick, Caitlin Kehoe, and Michael J Spivey. 2007. “Graded Motor Responses in the Time Course of Categorizing Atypical Exemplars.” *Memory & Cognition* 35 (1): 15–28.
- Davies, Catherine, Jamie Lingwood, Bissera Ivanova, and Sudha Arunachalam. 2021. “Three-Year-Olds’ Comprehension of Contrastive and Descriptive Adjectives: Evidence for Contrastive Inference.” *Cognition* 212 (July): 104707. <https://doi.org/10.1016/j.cognition.2021.104707>.
- Degen, Judith, Robert D Hawkins, Caroline Graf, Elisa Kreiss, and Noah D Goodman. 2020. “When Redundancy Is Useful: A Bayesian Approach to ‘Overinformative’ Referring Expressions.” *Psychological Review* 127: 591–621.
- Degen, Judith, and Michael K. Tanenhaus. 2015. “Processing Scalar Implicature: A Constraint-Based Approach.” *Cognitive Science* 39 (4): 667–710.
- Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2018. “Bert: Pre-Training of Deep Bidirectional Transformers for Language Understanding.” *arXiv Preprint arXiv:1810.04805*.
- Diesendruck, Gil, D. Geoffrey Hall, and Susan A. Graham. 2006. “Children’s Use of Syntactic

- and Pragmatic Knowledge in the Interpretation of Novel Adjectives.” *Child Development* 77 (1): 16–30.
- Engelhardt, Paul E., Ş. Barış Demiral, and Fernanda Ferreira. 2011. “Over-Specified Referring Expressions Impair Comprehension: An ERP Study.” *Brain and Cognition*, Special Section: Aggression and peer victimization: Genetic, neurophysiological, and neuroendocrine considerations, 77 (2): 304–14. <https://doi.org/10.1016/j.bandc.2011.07.004>.
- Firth, John R. 1957. “A Synopsis of Linguistic Theory, 1930-1955.” *Studies in Linguistic Analysis*.
- Foppolo, Francesca, Maria Teresa Guasti, and Gennaro Chierchia. 2012. “Scalar Implicatures in Child Language: Give Children a Chance.” *Language Learning and Development* 8 (4): 365–94.
- Frank, Michael C, and Noah D Goodman. 2012. “Predicting Pragmatic Reasoning in Language Games.” *Science* 336 (6084): 998–98.
- . 2014. “Inferring Word Meanings by Assuming That Speakers Are Informative.” *Cognitive Psychology* 75: 80–96.
- Frank, Michael C, Noah D Goodman, and Joshua B Tenenbaum. 2009. “Using Speakers’ Referential Intentions to Model Early Cross-Situational Word Learning.” *Psychological Science* 20 (5): 578–85.
- Gelman, Susan A., and Ellen M. Markman. 1985. “Implicit Contrast in Adjectives Vs. Nouns: Implications for Word-Learning in Preschoolers*.” *Journal of Child Language* 12 (1): 125–43.
- Goldin-Meadow, Susan, Susan C Levine, Larry V Hedges, Janellen Huttenlocher, Stephen W Raudenbush, and Steven L Small. 2014. “New Evidence About Language and Cognitive Development Based on a Longitudinal Study: Hypotheses for Intervention.” *American Psychologist* 69 (6): 588.

- Grice, H Paul. 1975. "Logic and Conversation." *1975*, 41–58.
- Harris, Paul L, and Melissa A Koenig. 2006. "Trust in Testimony: How Children Learn About Science and Religion." *Child Development* 77 (3): 505–24.
- Horowitz, Alexandra C., and Michael C. Frank. 2016. "Children's Pragmatic Inferences as a Route for Learning About the World." *Child Development* 87 (3): 807–19.
- Huang, Yi Ting, and Jesse Snedeker. 2009. "Semantic Meaning and Pragmatic Interpretation in 5-Year-Olds: Evidence from Real-Time Spoken Language Comprehension." *Developmental Psychology* 45 (6): 1723–39.
- . 2008. "Use of Referential Context in Children's Language Processing." *Proceedings of the 30th Annual Meeting of the Cognitive Science Society*, January.
- Johns, Brendan T, and Michael N Jones. 2012. "Perceptual Inference Through Global Lexical Similarity." *Topics in Cognitive Science* 4 (1): 103–20.
- Kanwisher, Nancy, Roger P. Woods, Marco Iacoboni, and John C. Mazziotta. 1997. "A Locus in Human Extrastriate Cortex for Visual Shape Analysis." *Journal of Cognitive Neuroscience* 9 (1): 133–42.
- Kennedy, Christopher. 2007. "Vagueness and Grammar: The Semantics of Relative and Absolute Gradable Adjectives." *Linguistics and Philosophy* 30 (1): 1–45. <https://doi.org/10.1007/s10988-006-9008-0>.
- Landau, Barbara, Lila R Gleitman, and Barbara Landau. 2009. *Language and Experience: Evidence from the Blind Child*. Vol. 8. Harvard University Press.
- Landauer, Thomas K, and Susan T Dumais. 1997. "A Solution to Plato's Problem: The Latent Semantic Analysis Theory of Acquisition, Induction, and Representation of Knowledge." *Psychological Review* 104 (2): 211.
- Legare, Cristine H, and Paul L Harris. 2016. "The Ontogeny of Cultural Learning." *Child Development* 87 (3): 633–42.
- Lewis, Molly, Martin Zettersten, and Gary Lupyan. 2019. "Distributional Semantics as a

- Source of Visual Knowledge.” *Proceedings of the National Academy of Sciences* 116 (39): 19237–8.
- Mangold, Roland, and Rupert Pobel. 1988. “Informativeness and Instrumentality in Referential Communication.” *Journal of Language and Social Psychology* 7 (3-4): 181–91.
- Mikolov, Tomas, Edouard Grave, Piotr Bojanowski, Christian Puhersch, and Armand Joulin. 2018. “Advances in Pre-Training Distributed Word Representations.” In *Proceedings of the International Conference on Language Resources and Evaluation (Lrec 2018)*.
- Mikolov, Tomas, Ilya Sutskever, Kai Chen, Greg S Corrado, and Jeff Dean. 2013. “Distributed Representations of Words and Phrases and Their Compositionality.” In *Advances in Neural Information Processing Systems*, 3111–9.
- Mitchell, Margaret, Ehud Reiter, and Kees van Deemter. 2013. “Typicality and Object Reference,” 7.
- Ni, Weijia. 1996. “Sidestepping Garden Paths: Assessing the Contributions of Syntax, Semantics and Plausibility in Resolving Ambiguities.” *Language and Cognitive Processes* 11 (3): 283–334.
- Noveck, Ira A. 2001. “When Children Are More Logical Than Adults: Experimental Investigations of Scalar Implicature.” *Cognition* 78 (2): 165–88.
- Papafragou, Anna, and Julien Musolino. 2003. “Scalar Implicatures: Experiments at the Semantics–Pragmatics Interface.” *Cognition* 86 (3): 253–82.
- Pechmann, Thomas. 1989. “Incremental Speech Production and Referential Overspecification.” *Linguistics* 27 (1): 89–110.
- Rhodes, Marjorie, Sarah-Jane Leslie, and Christina M Tworek. 2012. “Cultural Transmission of Social Essentialism.” *Proceedings of the National Academy of Sciences* 109 (34): 13526–31.
- Rogers, Timothy T, and James L McClelland. 2004. *Semantic Cognition: A Parallel Distributed Processing Approach*. MIT press.

- Rosch, Eleanor, Carol Simpson, and R Scott Miller. 1976. "Structural Bases of Typicality Effects." *Journal of Experimental Psychology: Human Perception and Performance* 2 (4): 491.
- Rubio-Fernández, Paula. 2016. "How Redundant Are Redundant Color Adjectives? An Efficiency-Based Analysis of Color Overspecification." *Frontiers in Psychology* 7.
- Ryskin, Rachel, Chigusa Kurumada, and Sarah Brown-Schmidt. 2019. "Information Integration in Modulation of Pragmatic Inferences During Online Language Comprehension." *Cognitive Science* 43 (8): e12769.
- Řehůřek, Radim, and Petr Sojka. 2010. "Software Framework for Topic Modelling with Large Corpora." In *Proceedings of the LREC 2010 Workshop on New Challenges for NLP Frameworks*, 45–50. Valletta, Malta: ELRA.
- Sedivy, Julie C. 2003. "Pragmatic Versus Form-Based Accounts of Referential Contrast: Evidence for Effects of Informativity Expectations." *Journal of Psycholinguistic Research* 32 (1): 3–23.
- Sedivy, Julie C., Michael K. Tanenhaus, Craig G. Chambers, and Gregory N. Carlson. 1999. "Achieving Incremental Semantic Interpretation Through Contextual Representation." *Cognition* 71 (2): 109–47.
- Sloutsky, Vladimir M, and Anna V Fisher. 2004. "Induction and Categorization in Young Children: A Similarity-Based Model." *Journal of Experimental Psychology: General* 133 (2): 166.
- Snow, Catherine E. 1972. "Mothers' Speech to Children Learning Language." *Child Development*, 549–65.
- Sperber, Dan, and Deirdre Wilson. 1986. *Relevance: Communication and Cognition*. Vol. 142. Citeseer.
- Stahl, Aimee E, and Lisa Feigenson. 2015. "Observing the Unexpected Enhances Infants' Learning and Exploration." *Science* 348 (6230): 91–94.

- Tessler, Michael Henry, Polina Tsvilodub, Jesse Snedeker, and Roger P. Levy. 2020. “Informational Goals, Sentence Structure, and Comparison Class Inference.” *Proceedings of the Annual Conference of the Cognitive Science Society*, January. <https://par.nsf.gov/biblio/10159025-informational-goals-sentence-structure-comparison-classes-inference>.
- Westerbeek, Hans, Ruud Koolen, and Alfons Maes. 2015. “Stored Object Knowledge and the Production of Referring Expressions: The Case of Color Typicality.” *Frontiers in Psychology* 6.
- Willits, Jon A, Rachel Shirley Sussman, and Michael S Amato. 2008. “Event Knowledge Vs. Verb Knowledge.” In *Proceedings of the 30th Annual Conference of the Cognitive Science Society*, 2227–32.
- Xiang, Ming, Christopher Kennedy, Weijie Xu, and Timothy Leffel. 2022. “Pragmatic Reasoning and Semantic Convention: A Case Study on Gradable Adjectives.” *Semantics and Pragmatics* 15 (September): 9:EA–9:EA. <https://doi.org/10.3765/sp.15.9>.
- Xu, Fei, and Joshua B Tenenbaum. 2007. “Word Learning as Bayesian Inference.” *Psychological Review* 114 (2): 245.
- Yu, Chen, and Linda B Smith. 2007. “Rapid Word Learning Under Uncertainty via Cross-Situational Statistics.” *Psychological Science* 18 (5): 414–20.